



CrabNet hyperparameter benchmark: Pareto fronts from raw data (repeat-averaged objectives)


MAE vs RMSE (Spearman $\rho=0.97$)




MAE vs GPU runtime (Spearman $\rho=-0.57$)




MAE vs model size (Spearman $\rho=-0.05$)



RMSE vs GPU runtime (Spearman $\rho=-0.61$)



RMSE vs model size (Spearman $\rho=-0.07$)



GPU runtime vs model size (Spearman $\rho=0.37$)

